

Gradient Descent: The Foundation of ML Optimization

From Taylor Series to Modern Deep Learning

Nipun Batra and Teaching Staff

IIT Gandhinagar

August 28, 2025

Outline

The Core ML Problem

$$\min_{\theta} f(\theta)$$

The Core ML Problem

$$\min_{\theta} f(\theta)$$

Examples everywhere:

- Linear regression: $\min(y - \mathbf{X}\theta)^2$
- Neural networks: min classification/regression loss
- Logistic regression: min cross-entropy loss

The Core ML Problem

$$\min_{\theta} f(\theta)$$

Examples everywhere:

- Linear regression: $\min(y - \mathbf{X}\theta)^2$
- Neural networks: min classification/regression loss
- Logistic regression: min cross-entropy loss

Challenge: Most ML problems have no closed-form solution

Geometric Intuition: Climbing Mountains


Imagine hiking in fog to reach the valley:

- Feel slope beneath your feet
- **Strategy:** Step in steepest downhill direction
- **Gradient** = steepest uphill
- **Negative gradient** = steepest downhill

Geometric Intuition: Climbing Mountains

Imagine hiking in fog to reach the valley:

- Feel slope beneath your feet
- **Strategy:** Step in steepest downhill direction
- **Gradient** = steepest uphill
- **Negative gradient** = steepest downhill



Think!

Think!

Why does following $-\nabla f$ lead us toward the minimum?

Think!

Think!

Why does following $-\nabla f$ lead us toward the minimum?

(Solution in Appendix)

Why Taylor Series?

Key insight: If we can't solve $\min f(\mathbf{x})$ exactly, approximate $f(\mathbf{x})$ locally!

Why Taylor Series?

Key insight: If we can't solve $\min f(\mathbf{x})$ exactly, approximate $f(\mathbf{x})$ locally!

Univariate Taylor expansion:

$$f(x) \approx f(x_0) + f'(x_0)(x - x_0) + \frac{1}{2}f''(x_0)(x - x_0)^2$$

Why Taylor Series?


Key insight: If we can't solve $\min f(\mathbf{x})$ exactly, approximate $f(\mathbf{x})$ locally!

Univariate Taylor expansion:


$$f(x) \approx f(x_0) + f'(x_0)(x - x_0) + \frac{1}{2}f''(x_0)(x - x_0)^2$$

- **Zero-order:** $f(x) \approx f(x_0)$ (constant)
- **First-order:** adds linear term (tangent)
- **Second-order:** adds quadratic curvature

Visual: Tangent Line Approximation



Visual: Adding Quadratic Term



Key insight: Higher-order terms give better approximations!

Concrete Example: $f(x) = \cos(x)$ at $x_0 = 0$

$$f(0) = \cos(0) = 1 \quad (1)$$

$$f'(0) = -\sin(0) = 0 \quad (2)$$

$$f''(0) = -\cos(0) = -1 \quad (3)$$

$$f'''(0) = \sin(0) = 0 \quad (4)$$

$$f^{(4)}(0) = \cos(0) = 1 \quad (5)$$

Concrete Example: $f(x) = \cos(x)$ at $x_0 = 0$

$$f(0) = \cos(0) = 1 \quad (1)$$

$$f'(0) = -\sin(0) = 0 \quad (2)$$

$$f''(0) = -\cos(0) = -1 \quad (3)$$

$$f'''(0) = \sin(0) = 0 \quad (4)$$

$$f^{(4)}(0) = \cos(0) = 1 \quad (5)$$

Taylor approximations:

$$\text{0th order: } f(x) \approx 1 \quad (6)$$

$$\text{2nd order: } f(x) \approx 1 - \frac{x^2}{2} \quad (7)$$

$$\text{4th order: } f(x) \approx 1 - \frac{x^2}{2} + \frac{x^4}{24} \quad (8)$$

Multivariate Taylor Series

$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x} + \frac{1}{2} \Delta \mathbf{x}^T \nabla^2 f(\mathbf{x}_0) \Delta \mathbf{x}$$

Multivariate Taylor Series


$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x} + \frac{1}{2} \Delta \mathbf{x}^T \nabla^2 f(\mathbf{x}_0) \Delta \mathbf{x}$$

where $\Delta \mathbf{x} = \mathbf{x} - \mathbf{x}_0$


Multivariate Taylor Series

$$f(\mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x} + \frac{1}{2} \Delta \mathbf{x}^T \nabla^2 f(\mathbf{x}_0) \Delta \mathbf{x}$$

where $\Delta \mathbf{x} = \mathbf{x} - \mathbf{x}_0$



Visual: Multivariate Case with Level Sets



Key: Gradient \perp level sets, tangent plane \perp gradient

From Taylor Series to Gradient Descent

Goal: Find $\Delta \mathbf{x}$ such that $f(\mathbf{x}_0 + \Delta \mathbf{x}) < f(\mathbf{x}_0)$

From Taylor Series to Gradient Descent

Goal: Find $\Delta \mathbf{x}$ such that $f(\mathbf{x}_0 + \Delta \mathbf{x}) < f(\mathbf{x}_0)$

Using first-order Taylor approximation:

$$f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$$

From Taylor Series to Gradient Descent

Goal: Find $\Delta \mathbf{x}$ such that $f(\mathbf{x}_0 + \Delta \mathbf{x}) < f(\mathbf{x}_0)$

Using first-order Taylor approximation:

$$f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$$

To minimize, we need: $\nabla f(\mathbf{x}_0)^T \Delta \mathbf{x} < 0$

From Taylor Series to Gradient Descent

Goal: Find $\Delta \mathbf{x}$ such that $f(\mathbf{x}_0 + \Delta \mathbf{x}) < f(\mathbf{x}_0)$

Using first-order Taylor approximation:

$$f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$$

To minimize, we need: $\nabla f(\mathbf{x}_0)^T \Delta \mathbf{x} < 0$

Vector geometry insight: For vectors \mathbf{a} , \mathbf{b} : $\mathbf{a}^T \mathbf{b} = |\mathbf{a}| |\mathbf{b}| \cos(\theta)$

Minimum when $\cos(\theta) = -1$ (opposite directions!)

From Taylor Series to Gradient Descent

Goal: Find $\Delta \mathbf{x}$ such that $f(\mathbf{x}_0 + \Delta \mathbf{x}) < f(\mathbf{x}_0)$

Using first-order Taylor approximation:

$$f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$$

To minimize, we need: $\nabla f(\mathbf{x}_0)^T \Delta \mathbf{x} < 0$

Vector geometry insight: For vectors \mathbf{a} , \mathbf{b} : $\mathbf{a}^T \mathbf{b} = |\mathbf{a}| |\mathbf{b}| \cos(\theta)$

Minimum when $\cos(\theta) = -1$ (opposite directions!)

Optimal choice: $\Delta \mathbf{x} = -\alpha \nabla f(\mathbf{x}_0)$ where $\alpha > 0$

From Taylor Series to Gradient Descent

Goal: Find $\Delta \mathbf{x}$ such that $f(\mathbf{x}_0 + \Delta \mathbf{x}) < f(\mathbf{x}_0)$

Using first-order Taylor approximation:

$$f(\mathbf{x}_0 + \Delta \mathbf{x}) \approx f(\mathbf{x}_0) + \nabla f(\mathbf{x}_0)^T \Delta \mathbf{x}$$

To minimize, we need: $\nabla f(\mathbf{x}_0)^T \Delta \mathbf{x} < 0$

Vector geometry insight: For vectors \mathbf{a} , \mathbf{b} : $\mathbf{a}^T \mathbf{b} = |\mathbf{a}| |\mathbf{b}| \cos(\theta)$

Minimum when $\cos(\theta) = -1$ (opposite directions!)

Optimal choice: $\Delta \mathbf{x} = -\alpha \nabla f(\mathbf{x}_0)$ where $\alpha > 0$

$$\mathbf{x}_{\text{new}} = \mathbf{x}_{\text{old}} - \alpha \nabla f(\mathbf{x}_{\text{old}})$$

The Gradient Descent Algorithm

$$\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t)$$

The Gradient Descent Algorithm

$$\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t)$$

Algorithm:

- ① **Initialize:** θ_0 (random or educated guess)
- ② **For** $t = 0, 1, 2, \dots$ until convergence:
 - Compute gradient: $\mathbf{g}_t = \nabla f(\theta_t)$
 - Update parameters: $\theta_{t+1} = \theta_t - \alpha \mathbf{g}_t$
 - Check convergence: $|\mathbf{g}_t| < \epsilon$

The Gradient Descent Algorithm

$$\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t)$$

Algorithm:

- ① **Initialize:** θ_0 (random or educated guess)
- ② **For** $t = 0, 1, 2, \dots$ until convergence:
 - Compute gradient: $\mathbf{g}_t = \nabla f(\theta_t)$
 - Update parameters: $\theta_{t+1} = \theta_t - \alpha \mathbf{g}_t$
 - Check convergence: $|\mathbf{g}_t| < \epsilon$

Key properties:

- First-order method (uses gradients, not Hessians)
- Greedy local search
- Guaranteed convergence for convex functions

Linear Regression Problem

Learn: $y = \theta_0 + \theta_1 x$ from data

x	y
1	1
2	2
3	3

Linear Regression Problem

Learn: $y = \theta_0 + \theta_1 x$ from data

x	y
1	1
2	2
3	3

Cost Function (Mean Squared Error):

$$\text{MSE}(\theta_0, \theta_1) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

Linear Regression Problem

Learn: $y = \theta_0 + \theta_1 x$ from data

x	y
1	1
2	2
3	3

Cost Function (Mean Squared Error):

$$\text{MSE}(\theta_0, \theta_1) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)^2$$

Goal: $(\theta_0^*, \theta_1^*) = \text{argmin}_{\theta_0, \theta_1} \text{MSE}(\theta_0, \theta_1)$

Computing Gradients for Linear Regression

We need: $\nabla \text{MSE} = \begin{bmatrix} \frac{\partial \text{MSE}}{\partial \theta_0} \\ \frac{\partial \text{MSE}}{\partial \theta_1} \end{bmatrix}$

Computing Gradients for Linear Regression

We need: $\nabla \text{MSE} = \begin{bmatrix} \frac{\partial \text{MSE}}{\partial \theta_0} \\ \frac{\partial \text{MSE}}{\partial \theta_1} \end{bmatrix}$

Partial derivatives:

$$\frac{\partial \text{MSE}}{\partial \theta_0} = \frac{2}{n} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)(-1) = -\frac{2}{n} \sum_{i=1}^n \epsilon_i \quad (9)$$

$$\frac{\partial \text{MSE}}{\partial \theta_1} = \frac{2}{n} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)(-x_i) = -\frac{2}{n} \sum_{i=1}^n \epsilon_i x_i \quad (10)$$

where $\epsilon_i = y_i - \hat{y}_i$ is the residual.

Computing Gradients for Linear Regression

We need: $\nabla \text{MSE} = \begin{bmatrix} \frac{\partial \text{MSE}}{\partial \theta_0} \\ \frac{\partial \text{MSE}}{\partial \theta_1} \end{bmatrix}$

Partial derivatives:

$$\frac{\partial \text{MSE}}{\partial \theta_0} = \frac{2}{n} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)(-1) = -\frac{2}{n} \sum_{i=1}^n \epsilon_i \quad (9)$$

$$\frac{\partial \text{MSE}}{\partial \theta_1} = \frac{2}{n} \sum_{i=1}^n (y_i - \theta_0 - \theta_1 x_i)(-x_i) = -\frac{2}{n} \sum_{i=1}^n \epsilon_i x_i \quad (10)$$

where $\epsilon_i = y_i - \hat{y}_i$ is the residual.

Matrix form: $\nabla \text{MSE}(\theta) = \frac{2}{n} \mathbf{X}^T (\mathbf{X}\theta - \mathbf{y})$

Step-by-Step Example

Initial: $\theta_0 = 4$, $\theta_1 = 0$, $\alpha = 0.1$

Step-by-Step Example

Initial: $\theta_0 = 4$, $\theta_1 = 0$, $\alpha = 0.1$

Iteration 1:

- Predictions: $\hat{y}_1 = 4$, $\hat{y}_2 = 4$, $\hat{y}_3 = 4$
- Errors: $\epsilon_1 = -3$, $\epsilon_2 = -2$, $\epsilon_3 = -1$
- $\frac{\partial \text{MSE}}{\partial \theta_0} = -\frac{2}{3}(-6) = 4$
- $\frac{\partial \text{MSE}}{\partial \theta_1} = -\frac{2}{3}(-10) = 6.67$
- $\theta_0 = 4 - 0.1 \times 4 = 3.6$
- $\theta_1 = 0 - 0.1 \times 6.67 = -0.67$

Step-by-Step Example


Initial: $\theta_0 = 4$, $\theta_1 = 0$, $\alpha = 0.1$

Iteration 1:

- Predictions: $\hat{y}_1 = 4$, $\hat{y}_2 = 4$, $\hat{y}_3 = 4$
- Errors: $\epsilon_1 = -3$, $\epsilon_2 = -2$, $\epsilon_3 = -1$
- $\frac{\partial \text{MSE}}{\partial \theta_0} = -\frac{2}{3}(-6) = 4$
- $\frac{\partial \text{MSE}}{\partial \theta_1} = -\frac{2}{3}(-10) = 6.67$
- $\theta_0 = 4 - 0.1 \times 4 = 3.6$
- $\theta_1 = 0 - 0.1 \times 6.67 = -0.67$

New parameters: $(\theta_0, \theta_1) = (3.6, -0.67)$

Visual: GD Path on Loss Surface



Notice: Algorithm takes larger steps when gradient is large!

The Gradient Descent Family

Three main variants based on data usage per update:

Method	Data per update	Updates per epoch	Convergence
Batch GD	n (all)	1	Smooth
SGD	1	n	Noisy
Mini-batch GD	b (batch size)	n/b	Balanced

The Gradient Descent Family

Three main variants based on data usage per update:

Method	Data per update	Updates per epoch	Convergence
Batch GD	n (all)	1	Smooth
SGD	1	n	Noisy
Mini-batch GD	b (batch size)	n/b	Balanced

Trade-offs: Computational cost vs. convergence stability vs. memory

The Gradient Descent Family

Three main variants based on data usage per update:


Method	Data per update	Updates per epoch	Convergence
Batch GD	n (all)	1	Smooth
SGD	1	n	Noisy
Mini-batch GD	b (batch size)	n/b	Balanced

Trade-offs: Computational cost vs. convergence stability vs. memory

Modern ML: Mini-batch GD with batch sizes 32-256 is most common

- Good balance of stability and efficiency
- Enables parallel computation (GPUs love batches!)

Learning Rate Effects



Pop Quiz #1

For dataset with 1000 samples and batch size 50:

- 1 How many iterations per epoch for mini-batch GD?
- 2 If SGD takes 1000 epochs to converge, roughly how many epochs should mini-batch take?
- 3 Why might SGD be noisier than batch GD?

Pop Quiz #1

For dataset with 1000 samples and batch size 50:

- 1 How many iterations per epoch for mini-batch GD?
- 2 If SGD takes 1000 epochs to converge, roughly how many epochs should mini-batch take?
- 3 Why might SGD be noisier than batch GD?

(Solutions in Appendix)

Convergence Rates for Convex Functions

L -smooth convex: $\|\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})\| \leq L\|\mathbf{x} - \mathbf{y}\|$

With step size $\alpha \in (0, 1/L]$:

$$f(\mathbf{x}_t) - f(\mathbf{x}^*) \leq \frac{L\|\mathbf{x}_0 - \mathbf{x}^*\|^2}{2t}$$

Convergence Rates for Convex Functions

L -smooth convex: $\|\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})\| \leq L\|\mathbf{x} - \mathbf{y}\|$

With step size $\alpha \in (0, 1/L]$:

$$f(\mathbf{x}_t) - f(\mathbf{x}^*) \leq \frac{L\|\mathbf{x}_0 - \mathbf{x}^*\|^2}{2t}$$

Rate: $O(1/t)$ (sublinear)

Convergence Rates for Convex Functions

L -smooth convex: $\|\nabla f(\mathbf{x}) - \nabla f(\mathbf{y})\| \leq L\|\mathbf{x} - \mathbf{y}\|$

With step size $\alpha \in (0, 1/L]$:

$$f(\mathbf{x}_t) - f(\mathbf{x}^*) \leq \frac{L\|\mathbf{x}_0 - \mathbf{x}^*\|^2}{2t}$$

Rate: $O(1/t)$ (sublinear)

μ -strongly convex + L -smooth:

$$f(\mathbf{x}_t) - f(\mathbf{x}^*) \leq \left(1 - \frac{\mu}{L}\right)^t (f(\mathbf{x}_0) - f(\mathbf{x}^*))$$

Rate: Linear convergence! Condition number $\kappa = L/\mu$

SGD as Unbiased Estimator

True gradient: $\nabla L(\theta) = \frac{1}{n} \sum_{i=1}^n \nabla \ell(f(\mathbf{x}_i; \theta), y_i)$

SGD gradient estimate: $\nabla \tilde{L}(\theta) = \nabla \ell(f(\mathbf{x}; \theta), y)$, where (\mathbf{x}, y) is sampled uniformly from $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

SGD as Unbiased Estimator

True gradient: $\nabla L(\theta) = \frac{1}{n} \sum_{i=1}^n \nabla \ell(f(\mathbf{x}_i; \theta), y_i)$

SGD gradient estimate: $\nabla \tilde{L}(\theta) = \nabla \ell(f(\mathbf{x}; \theta), y)$, where (\mathbf{x}, y) is sampled uniformly from $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Unbiased estimator property:

$$\mathbb{E}[\nabla \tilde{L}(\theta)] = \nabla L(\theta)$$

SGD as Unbiased Estimator

True gradient: $\nabla L(\theta) = \frac{1}{n} \sum_{i=1}^n \nabla \ell(f(\mathbf{x}_i; \theta), y_i)$

SGD gradient estimate: $\nabla \tilde{L}(\theta) = \nabla \ell(f(\mathbf{x}; \theta), y)$, where (\mathbf{x}, y) is sampled uniformly from $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Unbiased estimator property:

$$\mathbb{E}[\nabla \tilde{L}(\theta)] = \nabla L(\theta)$$

Proof sketch:

$$\mathbb{E}[\nabla \tilde{L}(\theta)] = \sum_{i=1}^n \frac{1}{n} \nabla \ell(f(\mathbf{x}_i; \theta), y_i) = \nabla L(\theta)$$

Implication: Individual SGD steps might be “wrong”, but average in correct direction!

Why Unbiasedness Matters

Intuitive analogy: Asking random people for directions

- Individual answers might be slightly off
- But if no systematic bias, average direction is correct
- SGD does the same with gradient estimates!

Why Unbiasedness Matters

Intuitive analogy: Asking random people for directions

- Individual answers might be slightly off
- But if no systematic bias, average direction is correct
- SGD does the same with gradient estimates!

The noise can be beneficial:

- Helps escape local minima in non-convex problems
- Provides implicit regularization
- Enables online learning

GD vs Normal Equation

For linear regression:

Normal equation: $\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

- Time complexity: $O(d^2 n + d^3)$

GD vs Normal Equation

For linear regression:

Normal equation: $\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

- Time complexity: $O(d^2 n + d^3)$

Gradient descent: $\theta_{t+1} = \theta_t - \alpha \mathbf{X}^T (\mathbf{X} \theta_t - \mathbf{y})$

- Time complexity per iteration: $O(dn)$
- Total: $O(T \cdot dn)$ for T iterations

GD vs Normal Equation

For linear regression:

Normal equation: $\hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

- Time complexity: $O(d^2 n + d^3)$

Gradient descent: $\theta_{t+1} = \theta_t - \alpha \mathbf{X}^T (\mathbf{X} \theta_t - \mathbf{y})$

- Time complexity per iteration: $O(dn)$
- Total: $O(T \cdot dn)$ for T iterations

When to use which?

- Few features ($d < 1000$): Normal equation
- Many features ($d > 10000$): Gradient descent
- Non-linear models: Only gradient descent works

Beyond Basic Gradient Descent

Momentum: $\mathbf{v}_{t+1} = \beta \mathbf{v}_t + (1 - \beta) \mathbf{g}_t$, $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \alpha \mathbf{v}_{t+1}$

Adam: Combines momentum + adaptive learning rates

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \frac{\alpha}{\sqrt{\hat{\mathbf{v}}_t} + \epsilon} \hat{\mathbf{m}}_t$$

Defaults: $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$

Beyond Basic Gradient Descent

Momentum: $\mathbf{v}_{t+1} = \beta \mathbf{v}_t + (1 - \beta) \mathbf{g}_t$, $\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \alpha \mathbf{v}_{t+1}$

Adam: Combines momentum + adaptive learning rates

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \frac{\alpha}{\sqrt{\hat{\mathbf{v}}_t} + \epsilon} \hat{\mathbf{m}}_t$$

Defaults: $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\epsilon = 10^{-8}$

Why these improvements?

- Handle different parameter scales automatically
- Accelerate convergence in relevant directions
- Reduce oscillations in narrow valleys

Second-Order Methods

Newton's method: $\theta_{t+1} = \theta_t - \alpha[\nabla^2 f(\theta_t)]^{-1} \nabla f(\theta_t)$

Gauss-Newton: For least squares problems

L-BFGS: Quasi-Newton method (approximates Hessian)

Newton's method: $\theta_{t+1} = \theta_t - \alpha[\nabla^2 f(\theta_t)]^{-1} \nabla f(\theta_t)$

Gauss-Newton: For least squares problems

L-BFGS: Quasi-Newton method (approximates Hessian)

Trade-off: Faster convergence vs. $O(d^3)$ cost per iteration

Newton's method: $\theta_{t+1} = \theta_t - \alpha[\nabla^2 f(\theta_t)]^{-1} \nabla f(\theta_t)$

Gauss-Newton: For least squares problems

L-BFGS: Quasi-Newton method (approximates Hessian)

Trade-off: Faster convergence vs. $O(d^3)$ cost per iteration

Line search methods: Adaptive step size via Armijo condition

Every modern deep learning framework uses GD variants!

Key extensions:

- **Backpropagation:** Efficient gradient computation
- **Automatic differentiation:** PyTorch/TensorFlow handle gradients
- **GPU acceleration:** Parallel mini-batch gradients
- **Mixed precision:** 16-bit + 32-bit arithmetic

Common strategies:

- Grid search: $\alpha \in \{0.001, 0.01, 0.1, 1.0\}$
- Learning rate schedules: Start high, decay over time
- Adaptive methods: Let algorithm adjust automatically
- Learning rate finder: Gradually increase and watch loss

Common strategies:

- Grid search: $\alpha \in \{0.001, 0.01, 0.1, 1.0\}$
- Learning rate schedules: Start high, decay over time
- Adaptive methods: Let algorithm adjust automatically
- Learning rate finder: Gradually increase and watch loss

Warning signs:

- Loss exploding $\Rightarrow \alpha$ too high
- Very slow convergence $\Rightarrow \alpha$ too low
- Oscillating loss \Rightarrow Try smaller α or momentum

Other Practical Considerations

Feature scaling: Standardize features: $(x - \mu)/\sigma$

Convergence criteria:

- Gradient magnitude: $\|\nabla f(\boldsymbol{\theta})\| < \epsilon$
- Function change: $|f(\boldsymbol{\theta}_{t+1}) - f(\boldsymbol{\theta}_t)| < \epsilon$
- Maximum iterations: Simple upper bound

Other Practical Considerations

Feature scaling: Standardize features: $(x - \mu)/\sigma$

Convergence criteria:

- Gradient magnitude: $\|\nabla f(\boldsymbol{\theta})\| < \epsilon$
- Function change: $|f(\boldsymbol{\theta}_{t+1}) - f(\boldsymbol{\theta}_t)| < \epsilon$
- Maximum iterations: Simple upper bound

Common pitfalls:

- Poor initialization (use Xavier/He for neural networks)
- Poor feature scaling (different parameter scales)
- Not monitoring validation performance

Think!

Think!

For L -smooth function, why might $\alpha > 2/L$ cause divergence on a quadratic?

Think!

Think!

For L -smooth function, why might $\alpha > 2/L$ cause divergence on a quadratic?

(Solution in Appendix)

What We've Learned

Core journey:

- **Mathematical foundation:** Taylor series approximation
- **Key insight:** Follow $-\nabla f$ for steepest descent
- **Algorithm:** $\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t)$
- **Variants:** Batch, SGD, mini-batch (use mini-batch!)
- **Theory:** Convergence rates depend on function properties

What We've Learned

Core journey:

- **Mathematical foundation:** Taylor series approximation
- **Key insight:** Follow $-\nabla f$ for steepest descent
- **Algorithm:** $\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t)$
- **Variants:** Batch, SGD, mini-batch (use mini-batch!)
- **Theory:** Convergence rates depend on function properties

From theory to practice:

- Tune learning rates carefully
- Scale features properly
- Monitor diagnostics
- Consider advanced optimizers (Adam, momentum)

What We've Learned

Core journey:

- **Mathematical foundation:** Taylor series approximation
- **Key insight:** Follow $-\nabla f$ for steepest descent
- **Algorithm:** $\theta_{t+1} = \theta_t - \alpha \nabla f(\theta_t)$
- **Variants:** Batch, SGD, mini-batch (use mini-batch!)
- **Theory:** Convergence rates depend on function properties

From theory to practice:


- Tune learning rates carefully
- Scale features properly
- Monitor diagnostics
- Consider advanced optimizers (Adam, momentum)

Gradient descent powers modern machine learning!


For comprehensive mathematical analysis:

X	y
$\begin{matrix} \text{---} x_1^T \text{---} \\ \vdots \\ \text{---} x_N^T \text{---} \end{matrix}$	$\begin{matrix} y_1 \\ \\ \\ y_N \end{matrix}$


X	y	$\hat{y} = f(x, \theta)$
x_1^T	y_1	\hat{y}_1
\vdots		\vdots
x_N^T	y_N	\hat{y}_N



LOSS SURFACE OVER
 $6N^2$ EXAMPLES



LOSS SURFACE OVER
 $6N^{\text{th}}$ EXAMPLES



LOSS SURFACE OVER
 $6N^{\circ}$ EXAMPLES

Think! 1: Why does $-\nabla f$ lead toward minimum?

The gradient $\nabla f(\mathbf{x})$ points in direction of steepest *ascent*. To *descend* (minimize), we go in the opposite direction: $-\nabla f(\mathbf{x})$.

Think! 1: Why does $-\nabla f$ lead toward minimum?

The gradient $\nabla f(\mathbf{x})$ points in direction of steepest *ascent*. To *descend* (minimize), we go in the opposite direction: $-\nabla f(\mathbf{x})$.

Think! 2: Why might $\alpha > 2/L$ cause divergence?

For quadratic $f(x) = \frac{L}{2}x^2$, we have $f'(x) = Lx$. The update becomes:

$$x_{t+1} = x_t - \alpha L x_t = (1 - \alpha L)x_t.$$

If $\alpha L > 2$, then $|1 - \alpha L| > 1$, causing the sequence to diverge.

Pop Quiz #1: For 1000 samples, batch size 50:

- ① Mini-batch iterations per epoch: $1000/50 = 20$
- ② If SGD takes 1000 epochs, mini-batch might take ≈ 50 epochs (rough estimate)
- ③ SGD is noisier because it uses only 1 sample per update vs. all samples for batch GD

Pop Quiz #1: For 1000 samples, batch size 50:

- ① Mini-batch iterations per epoch: $1000/50 = 20$
- ② If SGD takes 1000 epochs, mini-batch might take ≈ 50 epochs (rough estimate)
- ③ SGD is noisier because it uses only 1 sample per update vs. all samples for batch GD

Additional insight: The noise in SGD can actually help escape local minima in non-convex optimization problems!

References & Further Reading

Essential references:

- Boyd & Vandenberghe: *Convex Optimization*
- Nocedal & Wright: *Numerical Optimization*
- Goodfellow et al.: *Deep Learning*, Chapters 4 & 8

References & Further Reading

Essential references:

- Boyd & Vandenberghe: *Convex Optimization*
- Nocedal & Wright: *Numerical Optimization*
- Goodfellow et al.: *Deep Learning*, Chapters 4 & 8

Interactive resources:

- Gradient descent visualizations online
- Implement from scratch in NumPy/PyTorch
- Experiment with different learning rates

References & Further Reading

Essential references:

- Boyd & Vandenberghe: *Convex Optimization*
- Nocedal & Wright: *Numerical Optimization*
- Goodfellow et al.: *Deep Learning*, Chapters 4 & 8

Interactive resources:

- Gradient descent visualizations online
- Implement from scratch in NumPy/PyTorch
- Experiment with different learning rates

Next lecture: Advanced Optimization Techniques
Practice: Implement GD for your favorite ML model!